

System Architecture for Investor Follow-Up Workflow

Purpose

This workflow is for Sago to help an investor re-engage with a startup when meaningful signals indicate a company's progression. It reflects the three main principles of:

Seamless Integration: workflow uses only Gmail and Slack, so no new UI needed

Hyper-Personalization: utilizes past interactions (e.g., notes/emails) to tailor outreach

True Agency: generates and executes outreach only with approval

The system monitors company-level signals, generates contextual outreach, and executes the send via Gmail, all from within an investor's existing workflow.

High Level Architecture

Data Ingestion Layer

Gmail Connector:

- Imports email threads, meeting notes, and metadata
- Extracts contact details, relevant snippets, commitments, timelines (e.g., "reach back in 6–12 months"), and last contact date.
- Prototype uses a simplified JSON import rather than full OAuth.

Signal Feeds:

- Unified interface for consuming event data.
- Prototype uses mocked JSON signals (funding, product updates).
- In production: Crunchbase, AngelList, GitHub, LinkedIn hiring data, etc.

Slack Connector:

- Sends notifications and receives Approve / Edit / Snooze actions.
- Can later support interactive workflows.

Profile and Context Score

Each contact is represented by a small context graph that includes:

- Founder identity + company
- Last interaction date
- Extracted notes + conversation summary
- Embedding vector for similarity
- Investor-specific tags (e.g., "too early")
- Prior commitments (e.g., "Planning to raise Series A in December 2025")

Storage should be lightweight - think SQLite or JSON for prototyping.

Signal Engine

This represents the system's core decision layer. It performs the following:

1. Normalizes incoming signals
2. Scores them using recency, significance, and relevance to investor
3. Sends a “Re-Engagement Opportunity) if score > threshold

The scoring function considers the following points:

- Event strength (e.g., funding > GitHub commits)
- Event recency
- Alignment with past conversations (e.g., “reach out after we hire MLOps head”)
- Change in company maturity

Rules will be transparent for early prototyping, but have room for ML ranking later.

Outreach Generator

With a profile + signal in hand, the workflow will produce an outreach focusing on:

- Summarizing the signal
- Referencing relevant past interactions
- One-click approval path

This can support LLM-backed template generation or deterministic templates.

Agency + Execution Layer

Investor approval can occur via Slack or a rule. Once that happens:

1. Outreach sent through email
2. Event is logged in audit trail
3. Profile timeline is updated

This reflects “true agency”, because the outreach can only go through with investor approval, but the execution is handled through Sago without need for other investor intervention.

Summary of End-to-End Flow

1. Investor meets founder, decides company is “too early”
2. Gmail thread imported and context is extracted
3. Signals flow in
4. Signal Engine identifies meaningful progression
5. Generator creates outreach
6. Slack message delivered to investor with Approve/Edit/Snooze
7. Upon approval, email sent via Gmail
8. Company profile updated

Architecture Diagram

**Gmail —→ Data Ingestion Layer —→ Profile & Context Store —→ Signal Engine —→
Outreach Generator —→ Execution & Agency —→ Gmail Send**

